

Pre ML Server Setup

Part 1: BIOS/UEFI Configuration: (Enter BIOS/UEFI settings before installing Ubuntu)

Step 1: Enter BIOS Setup:

1. Power on or reboot the server.
2. When prompted (usually within a few seconds), press **F9** to enter **System Utilities** (BIOS setup).
3. You may also see a prompt for iLO – skip that for now unless you want to configure remote access.

Step 2 : Enable Virtualization (CPU):

In BIOS settings:

1. Navigate to:
System Configuration > BIOS/Platform Configuration (RBSU) > Processor Options
2. Enable the following options:
 - a. **AMD-V / Intel VT-x**: Enable (this allows virtualization extensions)
 - b. **SVM Mode** (for AMD CPUs): Enable
 - c. **SR-IOV**: Enable (optional, for networking virtualization)

Step 3 : Enable IOMMU / PCIe Passthrough (for GPU Virtualization) :

Still under **Processor Options**:

1. Enable:
 - a. **IOMMU**: Enable (Required for PCIe passthrough)
 - b. **Memory Mapped I/O above 4GB**: Enable (important for large GPUs)
 - c. **ACS Support** (optional but helpful for passthrough isolation): Enable

Step 4 : GPU Options :

If the server has discrete GPUs:

1. Go to:

System Configuration > BIOS/Platform Configuration (RBSU) > Advanced Options > PCIe Options

2. Ensure:

- **PCIe ARI Support:** Enabled
- **Above 4G Decoding:** Enabled (critical for GPU memory mapping)

Step 5 : Boot Settings :

To allow boot from USB/DVD:

1. Go to:

System Configuration > BIOS/Platform Configuration (RBSU) > Boot Options

2. Ensure:

- **UEFI Boot Mode:** Enabled
- Set **USB or DVD** as first boot device.

Optional: Disable Secure Boot if needed under **Secure Boot** settings.

Part 2: Ubuntu 22.04 Installation:

Step 6: Create a Bootable USB with Ubuntu:

On another computer:

1. Download Ubuntu 22.04 ISO from ubuntu.com.

2. Use **Rufus** (Windows) or **balenaEtcher** (macOS/Linux) to create a bootable USB.

Step 7 : **Boot from USB**

1. Insert the bootable USB into the HPE server.
2. Power on and press **F11** to select **One-Time Boot Menu**.
3. Select the USB drive.
4. Ubuntu installer will start.

Step 8: Ubuntu Server Installation:

1. Select language and keyboard layout.
2. Choose "Install Ubuntu Server."
3. Configure:
 - Network (static or DHCP)
 - Disk Partitioning: Use entire disk (unless custom layout needed)
 - Profile (name, username, password)
4. Select optional snaps (you can skip).
5. Installation begins. Reboot after it's done.

Step 9: Post-Install Configuration:

After reboot:

1. Login with your user credentials.

2. Update packages:
 - `sudo apt update && sudo apt upgrade -y`

Part 3: Verify Virtualization Support:

Step 10: Check Virtualization:

- `egrep -c '(vmx|svm)' /proc/cpuinfo`
Output should be **1 or more** (means VT-x/SVM is available)
- To verify IOMMU:
 - `dmesg | grep -i iommu`

Technical Terms :

PCIe:

PCIe is a **high-speed serial computer expansion bus standard** that connects high-performance components—like GPUs, SSDs, and network cards—to the motherboard.

Think of it as:

A **data highway** between your CPU and critical hardware like the GPU or NVMe SSD.

Key Concepts:

1. Lanes (x1, x4, x8, x16)

- Each **lane** consists of two pairs of wires: one for sending data, one for receiving.
- **More lanes = more bandwidth.**

- For example:
 - x1 → 1 lane (used for WiFi cards)
 - x4 → 4 lanes (used for NVMe SSDs)
 - x16 → 16 lanes (used for GPUs)

2. Versions (v1.0 → v5.0 and beyond)

Each newer version increases data rate:

- PCIe 3.0 (commonly used): ~1 GB/s per lane
- PCIe 4.0: ~2 GB/s per lane
- PCIe 5.0: ~4 GB/s per lane

A **PCIe 4.0 x16 GPU** can transfer up to **32 GB/s!**

What connects via PCIe?

- GPUs (NVIDIA, AMD)
- NVMe SSDs
- Network Interface Cards (10GbE, 25GbE, etc.)
- RAID controllers
- Capture cards

Why is PCIe Important for Virtualization?

- For GPU Passthrough or SR-IOV:
 1. PCIe allows direct communication between a **VM and a physical device** (e.g., GPU).
 2. Technologies like **IOMMU** rely on PCIe to isolate and safely share devices.

Example:

Your GPU is plugged into a **PCIe x16 slot**. If you're doing GPU passthrough to a VM, you're allowing that VM to directly use the GPU via the PCIe interface.

BIOS Settings :

1. **AMD-V / Intel VT-x** (a.k.a. Virtualization Extensions)
 - **Purpose:** Allows the CPU to support running virtual machines with near-native performance.
 - **Technical term:** These are hardware-assisted virtualization features.
 - **Intel** calls it **VT-x**, and **AMD** calls it **AMD-V**.
 - **Why enable it:** Required by hypervisors like KVM, VirtualBox, VMware, etc.
2. **SVM Mode** (Secure Virtual Machine)
 - **Specific to AMD processors.**
 - **Purpose:** Enables AMD-V instructions at the BIOS level.
 - **Why enable it:** Without this, AMD CPUs can't run virtual machines.
3. **SR-IOV** (Single Root I/O Virtualization)
 - **Purpose:** Allows a physical PCIe device (like a NIC or GPU) to appear as multiple **virtual** devices to VMs.
 - **Technical use:** Essential in high-performance environments, especially in cloud and network virtualization.
 - **Why enable it:** Lets VMs directly access slices of physical devices for performance.
4. **IOMMU** (Input-Output Memory Management Unit)
 - **Intel** calls it **VT-d**, **AMD** calls it **IOMMU**.
 - **Purpose:** Allows direct access of PCIe devices (e.g., GPUs) to virtual machines.
 - **Why enable it:** Needed for **PCI passthrough** – giving a full GPU or NIC to a VM.
5. Memory-Mapped I/O above 4GB

- **Purpose:** Allows the system to assign memory address space to PCIe devices above 4 GB, which is necessary for 64-bit addressable devices like modern GPUs.
 - **Why enable it:** Prevents resource conflicts and allows full GPU memory access.
6. ACS (Access Control Services)
- **Purpose:** Helps isolate PCIe devices into separate IOMMU groups.
 - **Why it matters:** When doing GPU passthrough, ACS makes sure each device can be safely assigned to a different VM.
7. PCIe ARI (Alternative Routing-ID Interpretation)
- **Purpose:** Enhances how many PCIe functions a single physical slot can support.
 - **Why enable it:** Helps with SR-IOV and multi-function devices like GPUs with multiple display pipelines or NICs with virtual functions.
8. Above 4G Decoding
- **Purpose:** Allows the use of large address spaces (>4GB) for PCIe devices like GPUs.
 - **Why enable it:** Critical for systems with multiple GPUs or high-memory GPUs (e.g., 16GB, 24GB).
9. UEFI Boot Mode
- **UEFI** is the modern replacement for legacy BIOS.
 - **Why enable it:** Required by modern OSes like Ubuntu 22.04 for features like Secure Boot and faster booting.
10. **Secure Boot**
- **Purpose:** Prevents unauthorized OS or bootloaders from running.
 - **Why you might disable it:** If you're using unsigned drivers (e.g., NVIDIA proprietary driver), Secure Boot might block them.

General Server Statics Commands:

1. All PCI Devices (GPUs, NICs, etc.)

- `lspci`

Lists all devices connected to **PCIe/PCI** slots.

Add `-v` or `-vv` for more details:

- `lspci -nnv`

2. All USB Devices

`lsusb`

- Shows devices connected via USB ports.

3. Storage Devices (Disks, SSDs)

`lsblk`

- Lists block devices (like `/dev/sda`, `/dev/nvme0n1`).

`sudo fdisk -l`

- Detailed partition info.

4. Network Interfaces


```
ip link show
```

- Shows all network interfaces.

```
lshw -class network
```

- Detailed info about each network device.

5. Detailed Hardware Info (All Devices)

```
sudo lshw
```

- Full system hardware summary.
- To view a specific class, use:

```
sudo lshw -class display    # For GPU
```

```
sudo lshw -class processor  # For CPU
```

6. CPU & Memory Info

```
lscpu                      # Detailed CPU architecture
```

```
free -h          # Memory usage
```

7. IOMMU / PCI Grouping (for passthrough)

```
find /sys/kernel/iommu_groups/ -type l
```

- Lists IOMMU groups, helpful for GPU passthrough.
-

Bonus: View All Detected Devices

```
dmesg | less
```

- Kernel logs that include detected hardware during boot.